

SOC722 Exercise Week 1

Zhe Chen

2026-01-18

Exercise 1.1

Try typing `typeof(mtcars)` and `class(mtcars)` in the console to see what happens.

Now type the following chunks of code into your console and understand what they do:

- `nrow(mtcars)`
- `ncol(mtcars)`
- `length(mtcars)`
- `dim(mtcars)`
- `rownames(mtcars)`
- `colnames(mtcars)`

Briefly describe what each of these do.

```
typeof(mtcars)
```

[1] "list"

- The type of `mtcars` is `list`.

```
class(mtcars)
```

[1] "data.frame"

- The class of `mtcars` is `data.frame`.

```
nrow(mtcars)
```

[1] 32

- The number of rows in `mtcars` is 32.

```
ncol(mtcars)
```

[1] 11

- The number of columns in `mtcars` is 11.

```
length(mtcars)
```

[1] 11

- The length of `mtcars` is 11.

```
dim(mtcars)
```

[1] 32 11

- The dimension of `mtcars` is (32,11), i.e., it has 32 rows and 11 columns.

```
rownames(mtcars)
```

[1] "Mazda RX4"	"Mazda RX4 Wag"	"Datsun 710"
[4] "Hornet 4 Drive"	"Hornet Sportabout"	"Valiant"
[7] "Duster 360"	"Merc 240D"	"Merc 230"
[10] "Merc 280"	"Merc 280C"	"Merc 450SE"
[13] "Merc 450SL"	"Merc 450SLC"	"Cadillac Fleetwood"
[16] "Lincoln Continental"	"Chrysler Imperial"	"Fiat 128"
[19] "Honda Civic"	"Toyota Corolla"	"Toyota Corona"
[22] "Dodge Challenger"	"AMC Javelin"	"Camaro Z28"
[25] "Pontiac Firebird"	"Fiat X1-9"	"Porsche 914-2"
[28] "Lotus Europa"	"Ford Pantera L"	"Ferrari Dino"
[31] "Maserati Bora"	"Volvo 142E"	

- The names of rows (observations) in `mtcars` are “Mazda Rx4”, “Mazda Rx4 Wag”, “Datsun 710”...

```
colnames(mtcars)
```

```
[1] "mpg"   "cyl"   "disp"  "hp"    "drat"  "wt"    "qsec"  "vs"    "am"    "gear"  
[11] "carb"
```

- The names of columns (variables) in `mtcars` are “mpg”, “cyl”, “disp”...

Exercise 1.2

I suggest you always use long-form when creating logical vectors. Try assigning a different value to `TRUE` and to `T`.

Code

```
T <- 123  
TRUE <- 123
```

What just happened?

```
T <- 123  
T
```

```
[1] 123
```

```
TRUE <- 123
```

```
Error in TRUE <- 123: invalid (do_set) left-hand side to assignment
```

```
TRUE
```

```
[1] TRUE
```

- Assign value 123 to `T`. It succeeds.
- Assign value 123 to `TRUE`. It fails, and R returns `Error in TRUE <- 123 : invalid (do_set) left-hand side to assignment.`

Exercise 1.3

Implicit coercion

You can create atomic vectors of any length with `c()` for “concatenate”.

For example:

```
lg1 <- c(TRUE, FALSE, NA)
int <- c(1L, 6L, NA, 10L)
dbl <- c(1, NA, 2.5, 4.5)
chr <- c(NA, "these are", "some strings")
```

Recall that atomic vectors are homogeneous. If you try to concatenate vectors of different types you will end up discovering **implicit coercion**. Basically, different types will be coerced in the following order: logical → integer → double → character. For example, a logical and a character combine into a character:

```
str(c(TRUE, "chr")) ## str() is (almost) identical to dplyr::glimpse()
```

```
chr [1:2] "TRUE" "chr"
```

Test your knowledge of the vector coercion rules by predicting the output of the following uses of `c()`:

```
c(1, FALSE)
c("a", 1)
c(TRUE, 1L)
```

- `c(1, FALSE)` is a *double*.
- `c("a", 1)` is a *character*.
- `c(TRUE, 1L)` is an *integer*.

Exercise 1.4

Explicit coercion

Explicit coercion happens when you call a function like `as.logical()`, `as.integer()`, `as.double()`, or `as.character()`. Use `as.integer()` on FALSE and TRUE, what values do they get coerced to?

```
as.integer(FALSE)
```

```
[1] 0
```

- The value of `as.integer(FALSE)` is 0.

```
as.integer(TRUE)
```

```
[1] 1
```

- The value of `as.integer(TRUE)` is 1.

Exercise 1.5

The most common form of implicit coercion

The following chunk of code creates a logical vector of size 75.

```
x <- sample(c(TRUE, FALSE), size = 75, replace = TRUE)
str(x)
```

```
logi [1:75] FALSE FALSE TRUE TRUE TRUE TRUE ...
```

```
x <- sample(c(TRUE, FALSE), size = 75, replace = TRUE)
```

- The number of TRUE in `x` is:

```
sum(x)
```

```
[1] 39
```

- The proportion of TRUE in `x` is:

```
mean(x)
```

```
[1] 0.52
```

```
mean(x) == sum(x) / length(x)
```

```
[1] TRUE
```

- Yes, `mean(x)` and `sum(x) / length(x)` are equal.

Exercise 1.6

What is the difference between `mtcars["mpg"]` and `mtcars[["mpg"]]`? More generally, what is the difference between the `[` and `[[` operators?

Which of the following two is TRUE? “`identical(mtcars["mpg"], mtcars$mpg)`” `identical(mtcars[["mpg"]], mtcars$mpg)`”

```
mtcars["mpg"]
```

	mpg
Mazda RX4	21.0
Mazda RX4 Wag	21.0
Datsun 710	22.8
Hornet 4 Drive	21.4
Hornet Sportabout	18.7
Valiant	18.1
Duster 360	14.3
Merc 240D	24.4
Merc 230	22.8
Merc 280	19.2
Merc 280C	17.8
Merc 450SE	16.4
Merc 450SL	17.3
Merc 450SLC	15.2
Cadillac Fleetwood	10.4
Lincoln Continental	10.4
Chrysler Imperial	14.7
Fiat 128	32.4
Honda Civic	30.4
Toyota Corolla	33.9
Toyota Corona	21.5
Dodge Challenger	15.5
AMC Javelin	15.2
Camaro Z28	13.3
Pontiac Firebird	19.2
Fiat X1-9	27.3
Porsche 914-2	26.0
Lotus Europa	30.4
Ford Pantera L	15.8
Ferrari Dino	19.7
Maserati Bora	15.0
Volvo 142E	21.4

```
str(mtcars["mpg"])
```

```
'data.frame': 32 obs. of 1 variable:  
 $ mpg: num 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
```

```
mtcars[["mpg"]]
```

```
[1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8 16.4 17.3 15.2 10.4  
[16] 10.4 14.7 32.4 30.4 33.9 21.5 15.5 15.2 13.3 19.2 27.3 26.0 30.4 15.8 19.7  
[31] 15.0 21.4
```

```
str(mtcars[["mpg"]])
```

```
num [1:32] 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
```

- [] returns a data frame, while [[]] returns a vector (i.e., the content of mpg column in that data frame)

```
identical(mtcars["mpg"], mtcars$mpg)
```

```
[1] FALSE
```

```
identical(mtcars[["mpg"]], mtcars$mpg)
```

```
[1] TRUE
```

- identical(mtcars[["mpg"]], mtcars\$mpg) is TRUE.

Exercise 1.7

`letters` is a built-in object in R that contains the 26 letters of English alphabet.

Using the [operator, do the following: - Extract the 17th value of `letters` - Create a sequence of even numbers from 2 to 26 and use that to subset `letters` - Use 8:12 to subset `letters`.

This is known as integer subsetting.

What happens if instead of [you use [[:

```
letters
```

```
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s"  
[20] "t" "u" "v" "w" "x" "y" "z"
```

```
letters[17]
```

```
[1] "q"
```

```
even_numbers <- seq(2, 26, 2)  
letters[even_numbers]
```

```
[1] "b" "d" "f" "h" "j" "l" "n" "p" "r" "t" "v" "x" "z"
```

```
letters[8:12]
```

```
[1] "h" "i" "j" "k" "l"
```

```
letters[[8:12]]
```

```
Error in letters[[8:12]]: attempt to select more than one element in vectorIndex
```

- If using `[[8:12]]`, R returns `Error in letters[[8:12]] : attempt to select more than one element in vectorIndex`. Because `[]` only allows for extracting one element.

Exercise 1.8

Now that you know all this

Replace the 18th value of `letters` with a missing value (`NA`).

```
letters[18] <- NA  
letters
```

```
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" NA "s"  
[20] "t" "u" "v" "w" "x" "y" "z"
```

Exercise 1.9

Subset `mtcars` so that we only see the observations for which `cyl == 4`.

Subset `mtcars` so that we only see the observations for which `mpg` is greater than 23.

```
mtcars[mtcars$cyl == 4, ]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

```
mtcars[mtcars$mpg > 23, ]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2

Exercise 1.10

Using what I told you earlier about the `search()` function, explain why you get two different errors. What is going on? What is R doing when you type `table(year)`? (You might want to type `search()` into the console again). In what package does R find the `year` object?

```
table(year)
```

Error: object 'year' not found

```
search()
```

```
[1] ".GlobalEnv"      "package:stats"    "package:graphics"  
[4] "package:grDevices" "package:utils"     "package:datasets"  
[7] "package:methods"   "Autoloads"       "package:base"
```

- R returns Error: object 'year' not found. It is because neither `palmerpenguins` or `tidyverse` is loaded, and R cannot find the object `year` in the global environment.

```
library(tidyverse)
```

Warning: package 'purrr' was built under R version 4.5.2

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --  
v dplyr     1.1.4     v readr     2.1.5  
vforcats    1.0.0     v stringr   1.5.1  
v ggplot2   3.5.2     v tibble    3.3.0  
v lubridate 1.9.4     v tidyverse 1.3.1  
v purrr     1.2.0  
-- Conflicts ----- tidyverse_conflicts() --  
x dplyr::filter() masks stats::filter()  
x dplyr::lag()   masks stats::lag()  
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
library(palmerpenguins)
```

Attaching package: 'palmerpenguins'

The following objects are masked from 'package:datasets':

penguins, penguins_raw

```
table(year)
```

```
Error in unique.default(x, nmax = nmax): unique() applies only to vectors
```

```
search()
```

```
[1] ".GlobalEnv"           "package:palmerpenguins" "package:lubridate"  
[4] "package:forcats"       "package:stringr"        "package:dplyr"  
[7] "package:purrr"         "package:readr"          "package:tidyverse"  
[10] "package:tibble"        "package:ggplot2"        "package:tidyverse"  
[13] "package:stats"         "package:graphics"      "package:grDevices"  
[16] "package:utils"         "package:datasets"      "package:methods"  
[19] "Autoloads"            "package:base"
```

```
find("year")
```

```
[1] "package:lubridate"
```

```
year
```

```
function (x)
{
  UseMethod("year")
}
<bytecode: 0x0000028848ebe120>
<environment: namespace:lubridate>
```

- R returns `Error in unique.default(x, nmax = nmax): unique() applies only to vectors`. It is because R finds `year` in `tidyverse` which is a function, and R cannot apply `unique()` to a function.

Exercise 1.11

Use `slice()` to extract the even-numbered rows in the penguins dataset.

It will look something like this:

```
penguins |>
  slice("SOME NUMERIC VECTOR GOES HERE")
```

```
nrow(penguins)
```

```
[1] 344
```

```
penguins |>  
  slice(seq(2, 344, 2))
```

```
# A tibble: 172 x 8  
  species island   bill_length_mm bill_depth_mm flipper_length_mm body_mass_g  
  <fct>   <fct>        <dbl>        <dbl>          <dbl>        <int>  
1 Adelie  Torgersen     39.5       17.4           186        3800  
2 Adelie  Torgersen     NA          NA            NA          NA  
3 Adelie  Torgersen     39.3       20.6           190        3650  
4 Adelie  Torgersen     39.2       19.6           195        4675  
5 Adelie  Torgersen     42          20.2           190        4250  
6 Adelie  Torgersen     37.8       17.3           180        3700  
7 Adelie  Torgersen     38.6       21.2           191        3800  
8 Adelie  Torgersen     36.6       17.8           185        3700  
9 Adelie  Torgersen     42.5       20.7           197        4500  
10 Adelie Torgersen     46          21.5           194        4200  
# i 162 more rows  
# i 2 more variables: sex <fct>, year <int>
```

Now use `slice()` to extract every third row—i.e., row 3, 6, 9, and so on.

```
penguins |>  
  slice(seq(3, 344, 3))
```

```
# A tibble: 114 x 8  
  species island   bill_length_mm bill_depth_mm flipper_length_mm body_mass_g  
  <fct>   <fct>        <dbl>        <dbl>          <dbl>        <int>  
1 Adelie  Torgersen     40.3       18             195        3250  
2 Adelie  Torgersen     39.3       20.6           190        3650  
3 Adelie  Torgersen     34.1       18.1           193        3475  
4 Adelie  Torgersen     37.8       17.3           180        3700  
5 Adelie  Torgersen     34.6       21.1           198        4400  
6 Adelie  Torgersen     42.5       20.7           197        4500  
7 Adelie  Biscoe        37.8       18.3           174        3400  
8 Adelie  Biscoe        38.2       18.1           185        3950  
9 Adelie  Biscoe        40.6       18.6           183        3550
```

```

10 Adelie Biscoe           40.5      18.9       180      3950
# i 104 more rows
# i 2 more variables: sex <fct>, year <int>

```

Exercise 1.12

Use `filter()` to extract the observations in the penguins dataset for which `species == "Gentoo"`, `island == "Biscoe"`, and `body_mass_g` is between 5,000 and 5,500.

```

penguins |>
  filter(species == "Gentoo", island == "Biscoe") |>
  filter(body_mass_g >= 5000, body_mass_g <= 5500)

# A tibble: 39 x 8
  species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
  <fct>   <fct>        <dbl>        <dbl>          <int>       <int>
1 Gentoo  Biscoe       47.6        14.5          215       5400
2 Gentoo  Biscoe       46.7        15.3          219       5200
3 Gentoo  Biscoe       46.8        15.4          215       5150
4 Gentoo  Biscoe       48.7        15.1          222       5350
5 Gentoo  Biscoe       45.1        14.5          215       5000
6 Gentoo  Biscoe       46.3        15.8          215       5050
7 Gentoo  Biscoe       42.9        13.1          215       5000
8 Gentoo  Biscoe       46.1        15.1          215       5100
9 Gentoo  Biscoe       47.3        15.3          222       5250
10 Gentoo Biscoe       45.1        14.5          207       5050
# i 29 more rows
# i 2 more variables: sex <fct>, year <int>

```